

HH_Census_datasets_DataArts

June 25, 2020

In this notebook, we document and give a high level description of the Organizational, Household and Census data we have collected in our databases. Accessing this data require an userid and a password. The databases are hosted on a SQL server. Connecting to the server through an API using for example, python, would require necessary odbc driver.

1 Database: OrgDB

Import the general libraries first and connect to the SQL server

```
[1]: import pyodbc
import numpy as np
import pandas as pd
import os,sys
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(font_scale=1.5)
%matplotlib inline
import datetime
```

```
[2]: #- Reading my userid and password from the environment variables
hhuid=os.environ['HHUID']
hhpwd=os.environ['HHPWD']
#- Now we establish a connection so that we will be able to perform a query
→using pandas read_sql module
#- in python one can use pyodbc or create an engine using sqlalchemy
driver='/usr/local/lib/libmsodbcsql.17.dylib' #- local odbc drive
server='129.119.63.219'
dbname='OrgDB'
port=1433
cnxnOrg=pyodbc.connect(driver=driver,server=server,database=dbname,uid=hhuid,\
                        pwd=hhpwd,port=port)
```

Check the tables in the database

```
[3]: cursor = cnxnOrg.cursor()
for row in cursor.tables(tableType='TABLE'):
    if row[1]=='dbo': #- avoiding system tables
        print(row[2])
```

HHOrgData
HHOrgStatic

These tables have already been cleaned out from raw form and integrated for the static and organization level variables. We will explore each of these tables below.

```
[4]: def load_data(cnxn,sqlquery):  
    """  
    cnxn: pyodbc.Connection object  
    sqlquery: sql query string  
    returns pandas dataframe from the sqlquery.  
    Use only for small databases if running from standalone node-- to make_┐  
    →efficient  
    need distributed architecture for larger databases  
    """  
    cursor=cnxn.cursor()  
    data=pd.read_sql(sqlquery,cursor.connection)  
    return data
```

For data description we will limit our queries to a few rows. If one expects to extract the full table, it may be slow with the above function. One may increase the data loading efficiency by some form of parallel processing.

1.0.1 HHOrgStatic

Static information for the Organizations that have reported household transactions

```
[5]: #- look at the schema  
for row in cursor.columns(table='HHOrgStatic'):  
    print(row[3],row[5])
```

NCARID float
OrgID bigint
ORGName varchar
ADDRESS varchar
CITY varchar
STATE varchar
ZIP float
ZIP9 varchar
STATENO float
County float
FTRACT float
CensusBlock float
CNTYNM varchar
CBSA float
LATITUDE float
LONGITUDE float
Active bigint
InactiveDate float

sec_no float

```
[6]: sqlquery='select * from HHOrgStatic'
hhIntDF=load_data(cnxnOrg,sqlquery)
hhIntDF.head()
```

```
[6]:
```

	NCARID	OrgID	ORGName	ADDRESS	\
0	154202.0	1516	Barter Theatre	PO Box 867	
1	150159.0	186	WaterTower Theatre	15650 Addison Rd	
2	162722.0	851	Front Porch Theatricals	112 Sewickley Ridge Cir	
3	146464.0	1083	Baum School of Art	510 W Linden St	
4	146462.0	1084	Lehigh Valley Arts Council	840 Hamilton St	

	CITY	STATE	ZIP	ZIP9	STATENO	County	FTRACT	\
0	ABINGDON	VA	24212.0	24212-0867	51.0	51191.0	5.119101e+10	
1	ADDISON	TX	75001.0	75001-3285	48.0	48113.0	4.811301e+10	
2	ALEPPO TWP	PA	15143.0	15143-8978	42.0	42003.0	4.200356e+10	
3	ALLENTOWN	PA	18101.0	18101-1416	42.0	42077.0	4.207701e+10	
4	ALLENTOWN	PA	18101.0	18101-2455	42.0	42077.0	4.207701e+10	

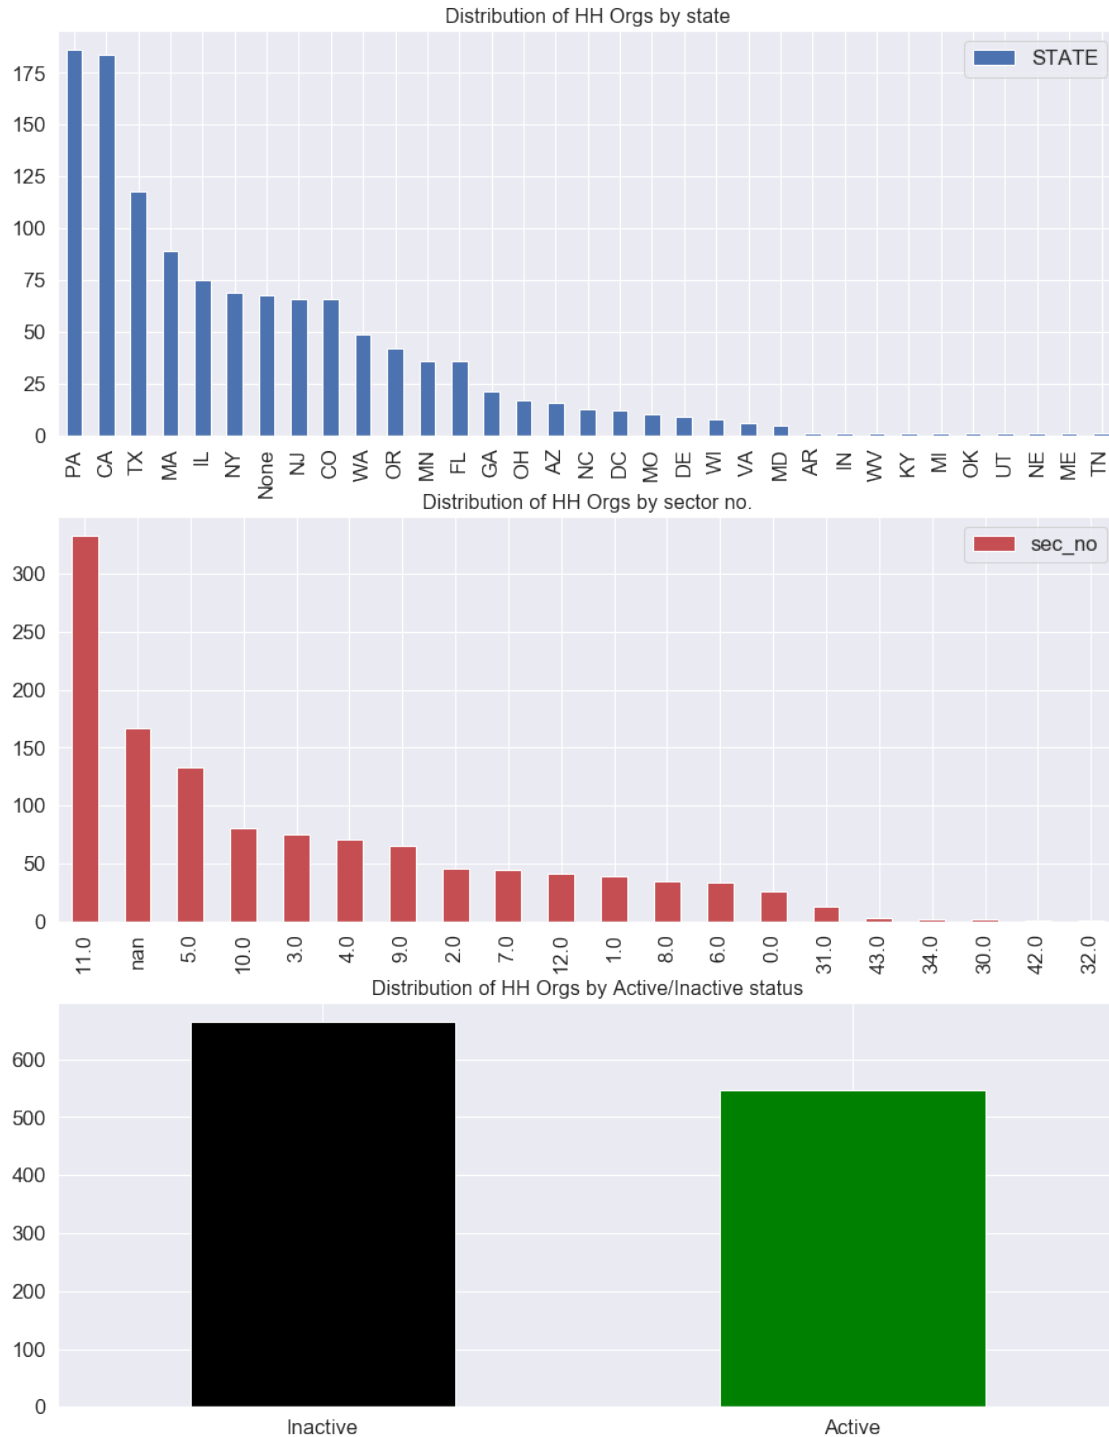
	CensusBlock	CNTYNM	CBSA	LATITUDE	LONGITUDE	Active	\
0	3011.0	WASHINGTON	28700.0	36.706928	-81.974386	0	
1	2012.0	DALLAS	19124.0	32.962209	-96.829781	1	
2	2028.0	ALLEGHENY	38300.0	40.540856	-80.147076	0	
3	1001.0	LEHIGH	10900.0	40.604335	-75.469017	0	
4	1035.0	LEHIGH	10900.0	40.601318	-75.474660	1	

	InactiveDate	sec_no
0	201807.0	11.0
1	NaN	11.0
2	201711.0	8.0
3	201709.0	1.0
4	NaN	3.0

```
[7]: #- Categorical Distributions
fig=plt.figure(figsize=(15,20))
ax1=plt.subplot(311)
hhIntDF['STATE'].astype(str).value_counts().plot(kind='bar')
ax1.legend()
plt.title('Distribution of HH Orgs by state',fontsize=16)
ax2=plt.subplot(312)
hhIntDF['sec_no'].astype(str).value_counts().plot(kind='bar',color='r')
ax2.legend()
plt.title('Distribution of HH Orgs by sector no.',fontsize=16)
ax3=plt.subplot(313)
hhIntDF['Active'].astype(str).value_counts().plot.bar(color=['Black','Green'])
```

```
#hhIntDF['Active'].astype(str).value_counts().
→plot(kind='bar',color=['Black','Green'],label='Inactive')
ax3.set_xticklabels(['Inactive','Active'],rotation=0)
plt.title('Distribution of HH Orgs by Active/Inactive status',fontsize=16)
```

[7]: Text(0.5, 1.0, 'Distribution of HH Orgs by Active/Inactive status')



1.0.2 HHOrgData

```
[8]: #- look at the schema
n=0
for row in cursor.columns(table='HHOrgData'):
    if n<=20:  #- only looking at the first 20 fields. Total 411
        if row[1]=='dbo':
            print(row[3],row[5])
        n+=1
```

OrgID bigint
year bigint
CNTART float
MKTADV float
ARTSATCD float
FRATNDTO float
PDATND float
ALLATTTO float
BOARDCD float
TRUSTNCD float
ENDTOTCD float
FTEMPS float
FTSEAS float
FTVOLS float
DEVSATCD float
GASAT float
HITIX float
LOTIX float
DMAILN float
MKTTOT float
MKTSAT float

HHOrgData Table consists of 411 variables with OrgID, year and the the remaining 409 numeric variables for the organizations (that have reported HH transactions) spanning from 2008 through 2019. Let's see some description below.

```
[9]: #Load the HH ORg data
sqlquery='select * from HHOrgData'
HHcompDF=load_data(cnxnOrg,sqlquery)
HHcompDF.head()
```

```
[9]:   OrgID  year  CNTART  MKTADV  ARTSATCD  FRATNDTO  PDATND  ALLATTTO  \
0   1012  2008  36598.0  7026.0         NaN    100.0    5003.0    5103.0
1   1012  2009  57887.0  8755.0         NaN    300.0    3925.0    4225.0
2   1012  2010  37799.0  2219.0         NaN     0.0  105260.0  105260.0
```

3	1012	2011	0.0	400.0	NaN	0.0	120545.0	120545.0
4	1012	2012	33819.0	713.0	76718.0	491.0	0.0	4100.0

	BOARD	CD	TRUST	NCD	...	GABEN	CD	PRGBEN	CD	UWEB	VIS	ArtsActivity	\
0	10.0		8.0		...	10268.0		21800.0		0.0		0.233247	
1	10.0		9.0		...	4325.0		25390.0		0.0		0.524995	
2	10.0		7.0		...	1662.0		22111.0		0.0		0.509109	
3	9.0		9.0		...	2088.0		26273.0		0.0		0.518762	
4	8.0		8.0		...	2983.0		24379.0		0.0		0.497554	

	ArtsProviders	GrantActivity	Hospitality	Substitute	SocioEcon	\
0	0.310376	-0.027368	0.187422	-0.147217	-0.093418	
1	0.191257	-0.017818	0.327811	-0.200109	-0.096941	
2	0.221737	0.444290	0.205557	-0.152727	0.383591	
3	0.341954	-0.444695	0.204866	-0.135336	0.431744	
4	0.184335	-0.444695	0.177981	-0.245605	0.528201	

	TOTPOP
0	11406.837973
1	11406.837973
2	11485.493201
3	11553.009420
4	11590.029118

[5 rows x 410 columns]

```
[10]: #- For display, let's take a subset and look at some correlation
selected_fields=['ArtsActivity', 'ArtsProviders',
                 'GrantActivity', 'Hospitality', 'Substitute', 'SocioEcon',
                 'TOTPOP']
HHcomp_subset=HHcompDF[selected_fields]
HHcomp_subset.describe()
```

[10]:	ArtsActivity	ArtsProviders	GrantActivity	Hospitality	Substitute	\
count	13246.000000	13246.000000	13246.000000	13246.000000	13246.000000	
mean	0.974061	2.271096	1.723004	1.094885	1.448059	
std	0.381326	1.519767	2.821947	0.798728	1.603717	
min	-1.655881	-1.041982	-0.444695	-1.328030	-0.879976	
25%	0.791436	1.200293	0.351602	0.504058	0.376510	
50%	0.994557	2.009976	1.047828	1.047564	1.275327	
75%	1.207564	2.867896	1.939133	1.691370	2.495204	
max	1.786953	6.376923	25.179070	3.550326	9.360109	

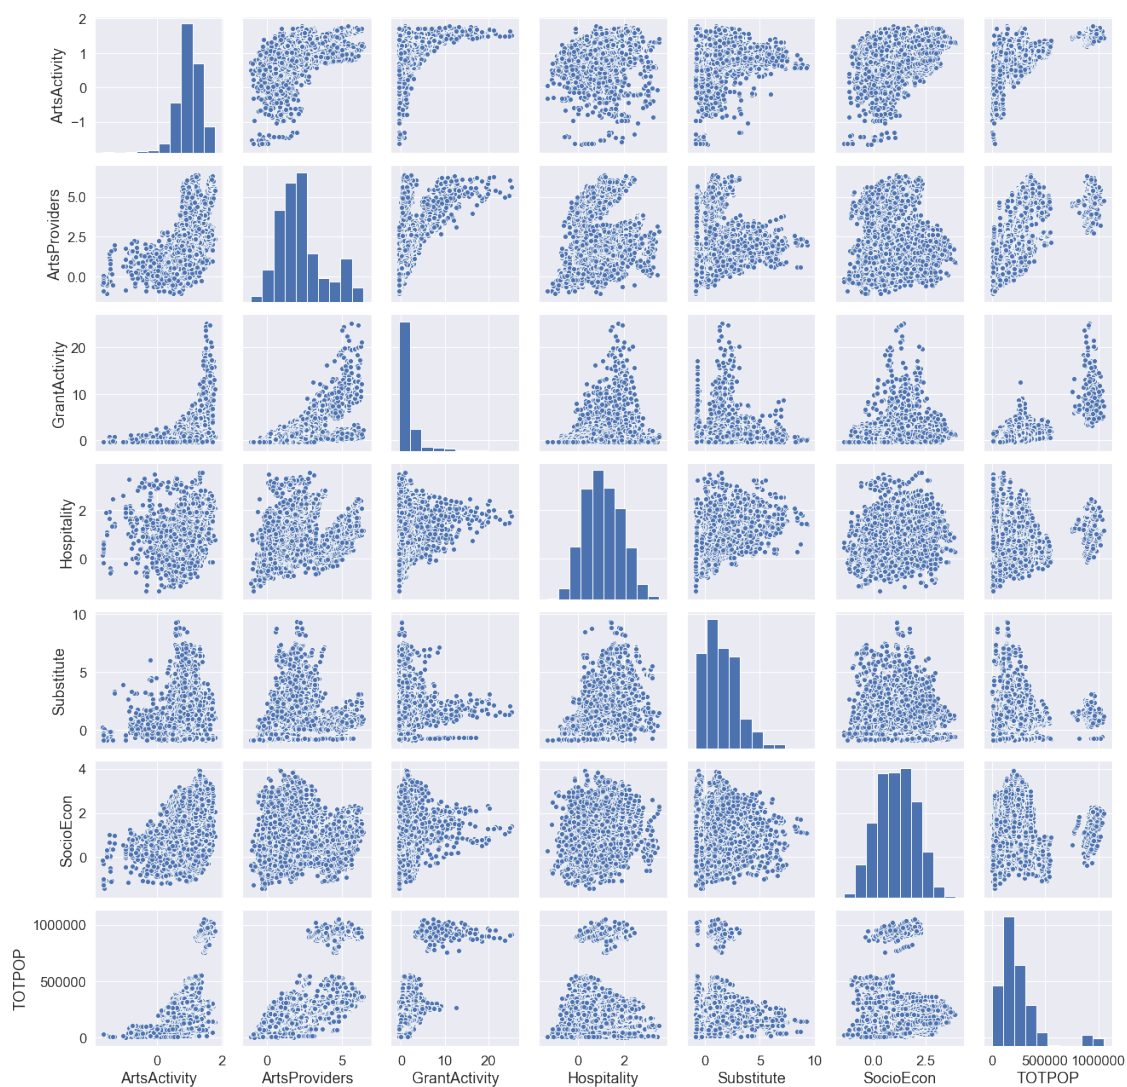
	SocioEcon	TOTPOP
count	13246.000000	1.324600e+04
mean	1.096317	2.470085e+05
std	0.900768	2.000732e+05

min	-1.422394	1.538331e+03
25%	0.420591	1.273388e+05
50%	1.082471	1.902962e+05
75%	1.760716	3.074388e+05
max	3.935416	1.051378e+06

One can look at the correlations in a pair plot

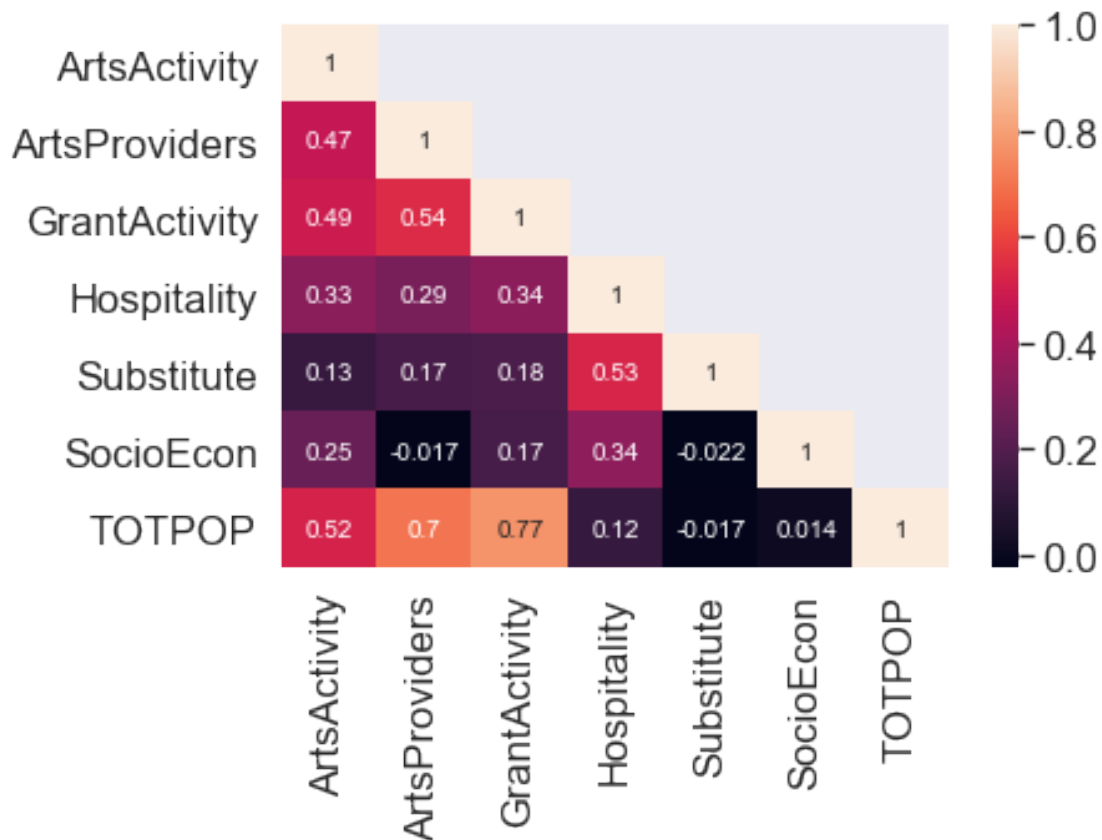
```
[11]: sns.pairplot(HHcomp_subset)
```

```
[11]: <seaborn.axisgrid.PairGrid at 0x1a236557d0>
```



Or one can also create a correlation matrix/see the overall correlation coefficients across the variables

```
[12]: corrMatrix=HHcomp_subset.corr()
corrMatrix=corrMatrix.where(np.tril(np.ones(corrMatrix.shape)).astype(np.bool))
#mask = np.triu(np.ones_like(corrMatrix, dtype=np.bool))
sns.heatmap(corrMatrix,annot=True)
plt.show()
```



The tables above can be joined by the ORGID/householdID. In this framework the join can be performed in the SQL query itself, or at the dataframe level. For larger tables, it is more efficient to perform the join operations in the SQL query itself

```
[13]: HHcomp_subset.head()
```

```
[13]:
```

	ArtsActivity	ArtsProviders	GrantActivity	Hospitality	Substitute	\
0	0.233247	0.310376	-0.027368	0.187422	-0.147217	
1	0.524995	0.191257	-0.017818	0.327811	-0.200109	
2	0.509109	0.221737	0.444290	0.205557	-0.152727	
3	0.518762	0.341954	-0.444695	0.204866	-0.135336	
4	0.497554	0.184335	-0.444695	0.177981	-0.245605	
	SocioEcon	TOTPOP				


```

0 -0.093418 11406.837973
1 -0.096941 11406.837973
2 0.383591 11485.493201
3 0.431744 11553.009420
4 0.528201 11590.029118

```

2 Database: HHDB

```

[14]: dbname='HHDB'
cnxnHH=pyodbc.connect(driver=driver,server=server,database=dbname,uid=hhuid,\
                        pwd=hhpwd,port=port)

```

Checking the tables in this DB

```

[15]: cursor = cnxnHH.cursor()
for row in cursor.tables(tableType='TABLE'):
    if row[1]=='dbo': #- avoiding system tables
        print(row[2])

```


HHActivity
HHStatic

2.0.1 Household Static

Static data showing geo coded Households

Distinct Households:

- **Total: 43,280,081**
- **State not NULL: 38,445,632**
- **US state+Territory: 38,048,817**



```

[16]: sqlquery='select top 100 * from HHStatic'
hshldDF=load_data(cnxnHH,sqlquery)
hshldDF.head()

```

```

[16]: HouseholdID CountyCode FTract BlockGroup City State PostalCode \
0 -40653585 None None None None None None None
1 -23727456 None None None None None None None
2 -139036295 None None None None None None None
3 -133529841 None None None Staten Island NY 10305
4 -124867765 None None None None None None None

```

	Fipsstatecode
0	NaN
1	NaN
2	NaN
3	36.0
4	NaN

2.0.2 Household Activity

```
[17]: #- look at the schema
for row in cursor.columns(table='HHActivity'):
    if row[1]=='dbo':
        print(row[3],row[5])
```



```
OrgID int
HouseholdID int
SegmentYear smallint
SegmentTypeCode varchar
SegmentDesc varchar
TransactionAmount money
TransactionQty int
OrderDate datetime
EventDate datetime
```

```
[18]: sqlquery='select top 100 * from HHActivity'
ActDF=load_data(cnxnHH,sqlquery)
ActDF.head()
```

```
[18]:
```

	OrgID	HouseholdID	SegmentYear	SegmentTypeCode	SegmentDesc	\
0	95	2480252	2014	GEN	Dabbler	
1	95	4166657	2014	GEN	Dabbler	
2	95	4290532	2014	GEN	Dabbler	
3	95	2571066	2014	GEN	Dabbler	
4	95	5990076	2014	GEN	Dabbler	

	TransactionAmount	TransactionQty	OrderDate	EventDate
0	None	None	None	None
1	None	None	None	None
2	None	None	None	None
3	None	None	None	None
4	None	None	None	None



The tables above can be joined by the ORGID/householdID. In this framework the join can be performed in the SQL query itself, or at the dataframe level. For larger tables, it is more efficient to perform the join operations in the SQL query itself

3 Database: CensusDB

We also have cleaned and integrated Census TRACT and Block Group level data that can be merged with the HH data for TRACT and Block group level analyses. For this, the database is CensusDB

```
[19]: #- We create a new connection instance
censusdb='CensusDB'
```

```
cnxnCNS=pyodbc.  
↪connect(driver=driver,server=server,database=censusdb,uid=hhuid,pwd=hhpwd,port=port)
```

```
[20]: #- checking the tables  
cursor = cnxnCNS.cursor()  
for row in cursor.tables(tableType='TABLE'):  
    if row[1]=='dbo': #- avoiding system tables  
        print(row[2])
```

```
BlkGrpcommute  
BlkGrpecon  
BlkGrpeduc  
BlkGrplatin  
BlkGrpLvl  
BlkGrpmedhhinc  
BlkGrppoverty  
BlkGrprace  
Tractdemo  
Tractecon  
Tracteduc  
Tracthshld  
TractLvl
```

The table names indicate the kinds of data in each table. The BlkGrp data span 2013-2019 and tract level data span from 2008-2019. The integrated tables are BlkGrpLvl and TractLvl and all the others are intermediate. Therefore we will only explore the final integrated tables at the Census Block Group and Census Tract level

3.0.1 BlkGrpLvl

```
[21]: sqlquery='select top 100 * from BlkGrpLvl'  
BlkGrpDF=load_data(cnxnCNS,sqlquery)  
BlkGrpDF.head()
```

```
[21]:
```

	YEAR	STATE	BlkGrp	CommuteN	AvgCommute	TotHse	LT50P	\
0	2013	Alabama	10010201001	268	14.082090	205	34.146341	
1	2013	Alabama	10010201002	570	33.156140	411	49.635036	
2	2013	Alabama	10010202001	535	27.691589	439	50.569476	
3	2013	Alabama	10010202002	398	26.097990	394	62.436548	
4	2013	Alabama	10010204001	501	21.055888	416	29.807692	

	GT100P	GT125P	GT150P	...	GradPlusP	MedHInc	WHITP	\
0	39.512195	10.243902	0.000000	...	11.616162	72375	86.656201	
1	16.058394	9.245742	6.812652	...	10.574413	52788	87.446627	
2	18.451025	3.189066	0.000000	...	3.363519	46979	30.296457	
3	14.974619	4.314721	3.045685	...	11.500701	43438	36.728395	
4	29.326923	19.230769	11.057692	...	9.948980	69375	97.794118	

	BLCKP	AMINDP	ASIAP	HAWAP	TOTPOP	NotLat	Latin
0	13.343799	0.000000	0.000000	0.0	637	637	0
1	5.380017	0.853971	0.000000	0.0	1171	1171	0
2	62.039046	0.000000	6.290672	0.0	1383	1334	49
3	62.448560	0.000000	0.000000	0.0	972	970	2
4	0.000000	2.022059	0.000000	0.0	1088	1072	16

[5 rows x 24 columns]

So that shows the Block Group level economic demographic, commute time etc by year for each Block Group.

3.0.2 TractLvl

```
[22]: sqlquery='select top 20 * from TractLvl'
      TractDF=load_data(cnxnCNS,sqlquery)
      TractDF.head()
```

```
[22]:  YEAR      STATE      TRACT  POP16  LT50P  GT100P  GT150P  GT200P  \
0  2011  California  6037575401   3598   21.5    1.6    0.0    0.0
1  2011  California  6037575402   2334   19.3    0.9    0.0    0.0
2  2011  California  6037575500     37    NaN    NaN    NaN    NaN
3  2011  California  6037575801   1783   11.7    4.4    0.0    0.0
4  2011  California  6037575802   3510   22.2    0.9    2.2    1.4

      MedHInc  POVPERC  ...  MarSize  MalHseSize  FemHseSize  NonFamSize  TotFam  \
0  35270.0    36.0  ...    4.63    6.14    4.38    1.56    854
1  30900.0    26.4  ...    3.80    3.88    4.12    1.57    751
2     NaN    45.9  ...    0.00    0.00    0.00    0.00     0
3  32344.0    38.2  ...    4.96    3.50    4.54    1.63    522
4  32109.0    37.0  ...    4.59    3.37    4.41    1.43    982

      AvFamSize  MARKID18  MALKID18  FEMKID18  SameSex
0     4.39      295      95      242      0.0
1     3.63      273     111      172      2.4
2     0.00       0       0       0      0.0
3     4.31     231       7     147      0.1
4     3.99     279     42     264      1.5
```

[5 rows x 37 columns]

This shows Tract level data for education, demographics, economy etc.

4 ASIDE – Combining aka merging aka joining data sets

We show two ways to merge the data sets and pick Tract level census data to do so as an example

```
[23]: #- Tract level Census data. we pick three tables
Tracttables=['Tractdemo','Tractecon','Tracteduc']
for tab in Tracttables:
    print("Table schema for : ", tab)
    for row in cursor.columns(table=tab):
        print(row[3],row[5])
```

```
Table schema for : Tractdemo
TRACT bigint
TOTPOP bigint
WHIT bigint
BLCK bigint
AMIND bigint
ASIA bigint
HAWA bigint
LATIN bigint
YEAR bigint
STATE varchar
Table schema for : Tractecon
TRACT bigint
POP16 bigint
LT50P float
GT100P float
GT150P float
GT200P float
MEDHINC float
PovPerc float
YEAR bigint
STATE varchar
Table schema for : Tracteduc
TRACT bigint
POP25 bigint
BACHP float
GRADP float
BachPlusP float
YEAR bigint
STATE varchar
```

As we see, we have TRACT, YEAR, STATE in all Tract tables, so we will use these to join the tables.

Using SQL join query – fast

```
[24]: sqlquery='select a.TRACT,a.YEAR,a.STATE,TOTPOP,WHIT,BLCK,AMIND,ASIA,HAWA,LATIN,\
POP16,LT50P,GT100P,GT150P,GT200P,PovPerc,\
POP25,BACHP,GRADP,BachPlusP from Tractecon a \
full outer join Tracteduc b \
on a.TRACT=b.TRACT and a.YEAR=b.YEAR and a.STATE=b.STATE \'
```

```
full outer join Tractdemo c \
on a.TRACT=c.TRACT and a.YEAR=c.YEAR and a.STATE=c.STATE'
```

```
[25]: #%%timeit
#TRACT_dataDF=load_data(cnxnCNS,sqlquery)
#==> 2min 16s ± 12.5 s per loop (mean ± std. dev. of 7 runs, 1 loop each)
```

```
[26]: t1 = datetime.datetime.now()
TRACT_dataDF=load_data(cnxnCNS,sqlquery)
t2 = datetime.datetime.now()
print("Time taken to execute the query and load to DF [Seconds] ", (t2-t1).
→seconds)
```

Time taken to execute the query and load to DF [Seconds] 149

```
[27]: print(TRACT_dataDF.shape)
TRACT_dataDF.head()
```

(814013, 20)

```
[27]:
```

	TRACT	YEAR	STATE	TOTPOP	WHIT	BLCK	AMIND	ASIA	HAWA	LATIN	\
0	1001020100	2008	Alabama	1852.0	1552.0	291.0	67.0	0.0	0.0	15.0	
1	1001020100	2010	Alabama	1809.0	1516.0	330.0	77.0	0.0	0.0	15.0	
2	1001020100	2011	Alabama	1768.0	1560.0	223.0	107.0	4.0	0.0	0.0	
3	1001020100	2013	Alabama	1808.0	1650.0	170.0	57.0	14.0	0.0	0.0	
4	1001020100	2016	Alabama	2010.0	1737.0	298.0	6.0	17.0	21.0	53.0	

	POP16	LT50P	GT100P	GT150P	GT200P	PovPerc	POP25	BACHP	\
0	1396	14.7	18.021468	1.88031	5.956905	9.091817	1234.0	11.050633	
1	1392	14.7	21.500000	2.00000	7.000000	10.500000	1242.0	13.700000	
2	1398	17.2	21.300000	4.90000	5.800000	10.200000	1284.0	10.800000	
3	1404	13.1	24.200000	4.90000	1.300000	10.500000	1162.0	15.700000	
4	1580	7.9	18.700000	8.10000	0.700000	9.900000	1298.0	16.600000	

	GRADP	BachPlusP
0	9.729163	20.750948
1	11.800000	25.400000
2	9.100000	19.900000
3	10.900000	26.700000
4	14.700000	31.400000

Using individual dataframe – slow

```
[28]: #%%timeit
#squery='select * from Tractecon'
#testDF=load_data(cnxnCNS,squery)
```

```
[29]: t3 = datetime.datetime.now()
      query1='select * from Tractecon'
      query2='select * from Tracteduc'
      query3='select * from Tractdemo'

      print("Reading Tract economy data")
      econDF=load_data(cnxnCNS,query1)
      print("Reading Tract education data")
      educDF=load_data(cnxnCNS,query2)
      print("Reading Tract demographics data")
      demoDF=load_data(cnxnCNS,query3)

      tract_mergeDF1=econDF.merge(educDF,on=['TRACT','YEAR','STATE'],how='outer')
      tract_mergeDF2=tract_mergeDF1.
      ↪merge(demoDF,on=['TRACT','YEAR','STATE'],how='outer')

      t4 = datetime.datetime.now()
      print("Time taken on full data queries and DF merge [Seconds] ", (t4-t3).
      ↪seconds)
```

Reading Tract economy data
 Reading Tract education data
 Reading Tract demographics data
 Time taken on full data queries and DF merge [Seconds] 152

```
[30]: print(tract_mergeDF2.shape)
      tract_mergeDF2.head()
```

(814013, 21)

```
[30]:
```

	TRACT	POP16	LT50P	GT100P	GT150P	GT200P	MEDHINC	\
0	1001020100	1396	14.7	18.021468	1.880310	5.956905	60255.0	
1	1001020200	1516	17.3	13.474851	0.298741	1.568570	34570.0	
2	1001020300	2549	21.8	11.497938	3.663303	0.458445	37101.0	
3	1001020400	3638	15.6	13.656101	3.482013	1.328458	48153.0	
4	1001020500	6948	12.5	18.500227	5.361798	0.970599	58256.0	

	PovPerc	YEAR	STATE	...	BACHP	GRADP	BachPlusP	TOTPOP	\
0	9.091817	2008	Alabama	...	11.050633	9.729163	20.750948	1852.0	
1	12.967858	2008	Alabama	...	14.157831	7.590959	21.930150	2045.0	
2	6.914586	2008	Alabama	...	11.327994	1.362692	12.643148	3443.0	
3	5.438941	2008	Alabama	...	13.756875	6.813766	20.713601	4639.0	
4	5.378651	2008	Alabama	...	21.315216	9.980556	31.834601	9339.0	

	WHIT	BLCK	AMIND	ASIA	HAWA	LATIN
0	1552.0	291.0	67.0	0.0	0.0	15.0
1	855.0	1128.0	0.0	22.0	0.0	6.0

2	2891.0	539.0	0.0	31.0	0.0	39.0
3	4486.0	85.0	22.0	14.0	0.0	128.0
4	8067.0	1131.0	88.0	146.0	0.0	471.0

[5 rows x 21 columns]